

# US Housing Market Prediction

BIG DATA - HANDS ON PROJECT  
USUC ALEXANDRU



# Introduction



- ❑ The housing market plays a crucial role in the US economy, affecting everything from personal wealth to national financial stability
- ❑ Predicting housing prices is a challenge due to the vast amount of data and the complex factors involved, such as location, economy, demographics and market trends.



# Objectives

The objectives for this project are:

- Analyze the large dataset with millions of records and over 3GB of real estate data
- Preprocess the data
- Generate helpful diagrams which will help us understand the dataset better
- Choose the relevant features for the training process





# Dataset

- The current dataset contains information from year 2012 until 2024, having features such as:
  - period of time
  - region
  - state code
  - property type
  - median sale price
  - median list price
  - median days on the market
  - mediana price per square foot
  - pending sales
  - homes sold, etc...
- Out of all 58 features that this dataset has, I chose just the relevant columns which will help the model predict the median sale price better.

region ▼	city ▼	state ▼	state_code ▼	property_type ▼	homes_sold ▼	homes_sold_mom ▼
Chicago, IL	Chicago	Illinois	IL	Multi-Family (2-4 Unit)	284.0	-0.2303523035230352
Parsippany, NJ	Parsippany	New Jersey	NJ	All Residential	14.0	0.75
Oakbrook, KY	Oakbrook	Kentucky	KY	All Residential	17.0	0.1333333333333333
Dunstable, MA	Dunstable	Massachusetts	MA	All Residential	6.0	0.0
Kalamazoo, MI	Kalamazoo	Michigan	MI	All Residential	42.0	-0.4473684210526315
Tysons, VA	Tysons	Virginia	VA	Condo/Co-op	11.0	-0.3529411764705882
Myrtle Creek, OR	Myrtle Creek	Oregon	OR	Single Family Residential	5.0	1.5
Valencia West, AZ	Valencia West	Arizona	AZ	Single Family Residential	25.0	0.3157894736842106
Erda, UT	Erda	Utah	UT	All Residential	3.0	-0.25

median_sale_price ▼	median_sale_price_mom ▼	median_sale_price_yoy ▼	median_list_price ▼	median_dom ▼	median_dom_mom ▼
259500.0	0.0176470588235293	0.1662921348314607	285000.0	42.0	6.0
485000.0	-0.0673076923076922	-0.0351801179467439	477500.0	33.0	3.0
265000.0	-0.0363636363636363	0.4887640449438202	260000.0	5.0	2.0
522500.0	0.1578947368421053	-0.0141509433962264	529900.0	50.0	41.0
187500.0	0.2798634812286688	0.25	179900.0	22.0	-1.0
325000.0	0.0483870967741935	0.0420006412311637	439450.0	23.0	-64.0
308000.0	0.3508771929824561	0.4322250639386189	209500.0	18.0	6.0
159000.0	-0.1740259740259739	0.0192307692307691	163386.0	72.0	-55.0
630000.0	0.0637220139803464	-0.0631970260223048	699900.0	163.0	37.0

# Tools

Tools used during the implementation of this project:

- Libraries
  - numpy
  - pandas
  - matplotlib
  - seaborn
  - sklearn
- Development environments
  - Pycharm
  - Jupyter notebook
  - Google colab

# Implementation

The implementation process consists of multiple steps:

1. Analyze the dataset
2. Cleansing
3. Add new relevant columns
4. Normalize the data
5. Generate helpful diagrams
6. Choose the most relevant columns for the training process
7. Train the model

# 1. Analyze the data

## 2. Cleansing

- drop records with null features

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 np.random.seed(0)
6
7 [1]
8
9 1 file_path = "dataset/us_city_market_tracker.csv"
10 2
11 3 us_ds = pd.read_csv(file_path, sep='\t')
12 4
13 [2]
14
15 1 us_ds = us_ds.dropna()
16 2
17 [4]
```

median_sale_price ▾ ↕	median_sale_price_mom ▾ ↕
35000	<null>
123145	0.21986131748390303
37478	-0.65616513761467887
142000	<null>
166945	<null>
121500	0.59474979491386382
949000	-0.66046511627906979
<null>	<null>
475000	-0.0083507306889353261
177500	-0.44357366771159878
119900	<null>



### 3. Add new relevant columns

```
1 us_ds['sale_to_list_ratio'] = us_ds['median_sale_price'] / us_ds['median_list_price']
2 us_ds['year'] = us_ds['period_begin'].dt.year
3 us_ds['month'] = us_ds['period_begin'].dt.month
4 us_ds['sales_to_new_listings_ratio'] = us_ds['homes_sold'] / us_ds['new_listings']
5 us_ds['price_increased_mom'] = (us_ds['median_sale_price_mom'] > 0).astype(int)
6 us_ds['price_increased_yoy'] = (us_ds['median_sale_price_yoy'] > 0).astype(int)
7 us_ds['inventory_turnover'] = us_ds['homes_sold'] / us_ds['inventory']
8 us_ds['sale_to_list_ppsf_ratio'] = us_ds['median_ppsf'] / us_ds['median_list_ppsf']
9 us_ds['supply_demand_balance'] = us_ds['new_listings'] - us_ds['pending_sales']
10 us_ds['fast_selling'] = (us_ds['median_dom'] <= 30).astype(int)

[7]
```

price_increased_yoy	inventory_turnover	sale_to_list_ppsf_ratio	supply_demand_balance	fast_selling
1	0.21580547112462006	0.517396987426864	55.0	0
0	0.7	0.9756112590691002	3.0	0
1	1.7	0.9250675723341937	3.0	1
0	0.4	1.1586413839521625	0.0	0
1	0.4158415841584158	0.9964415983172793	-8.0	1
1	0.12359550561797752	1.0343775541330438	16.0	1
1	0.45454545454545453	0.9206754735792623	1.0	1
1	0.26881720430107525	0.9594333074000767	-8.0	0
0	0.17647058823529413	0.9560499441946902	2.0	0
1	0.4308300395256917	0.9447901798458463	-3.0	1
1	3.8	0.9025138623310607	-2.0	1
1	0.5818181818181818	0.934585179102058	10.0	0
0	0.3	0.8344535203403132	0.0	1

## 4. Identify the outliers and replace them with moving average (MA)

```
4 exclude_columns = ['period_duration', 'region_type_id', 'table_id', 'property_type_id',
5 'parent_metro_region_metro_code', 'year', 'month', 'price_increased_mom',
6 'price_increased_yoy', 'fast_selling']
7
8 # us_ds_ca = us_ds[us_ds['state_code'] == 'CA']
9 numeric_columns = us_ds.select_dtypes(include=['number']).columns
10
11 if exclude_columns:
12     columns_for_outliers = [col for col in numeric_columns if col not in exclude_columns]
13 else:
14     columns_for_outliers = numeric_columns
15
16 z_scores = us_ds[columns_for_outliers].apply(zscore)
17
18 # define a threshold for outliers
19 threshold = 5
20
21 # identify outliers
22 outliers = (z_scores.abs() > threshold)
23
```

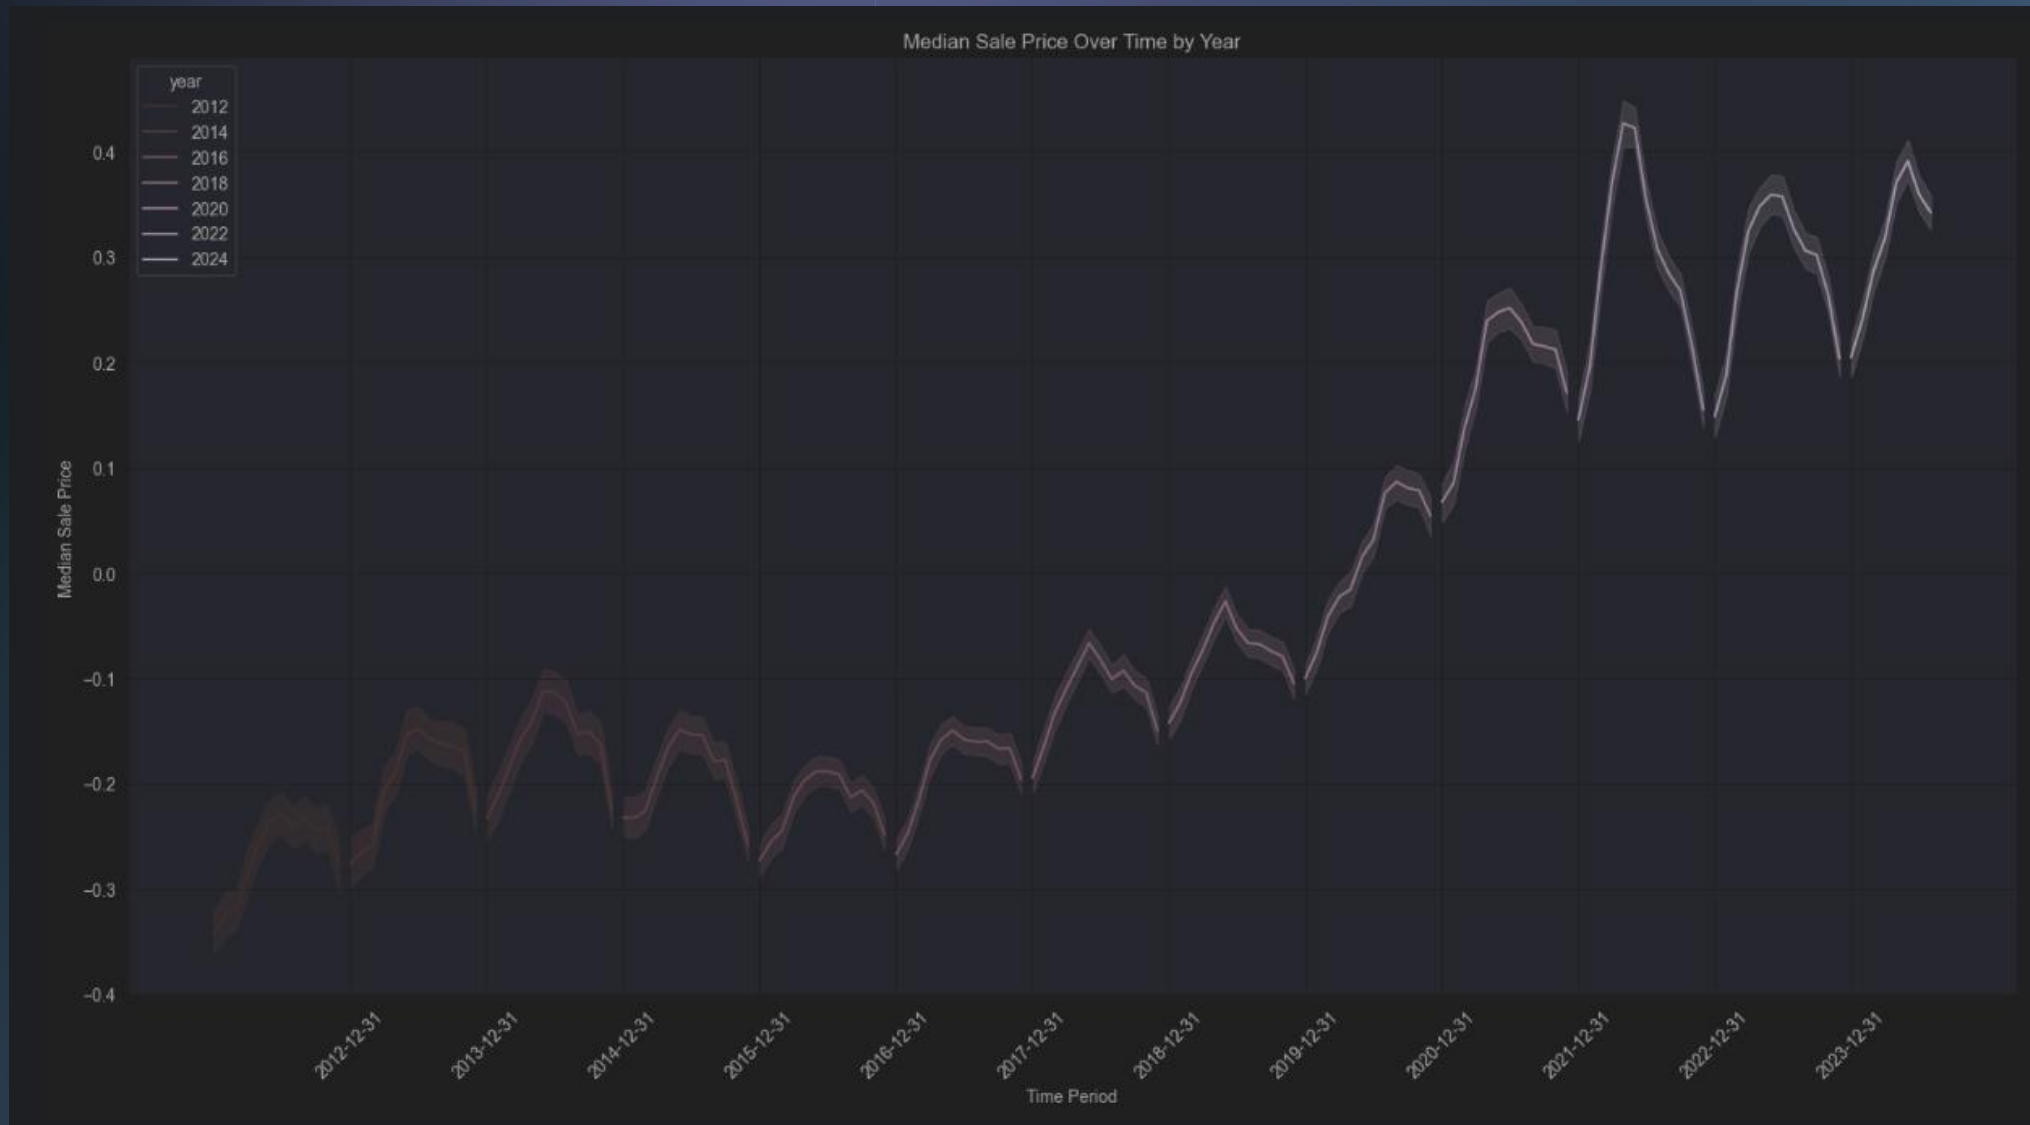
```
4 for col in columns_for_outliers:
5     # Calculate EMA for the column
6     # ema = us_ds_ca_copy[col].ewm(span=10, adjust=False).mean()
7
8     # calculate MA - moving average
9     ma = us_ds_copy[col].rolling(window=2, center=True).mean()
10
11     # Replace outliers with EMA
12     # us_ds_ca_copy.loc[outliers[col], col] = ema[outliers[col]]
13     us_ds_copy.loc[outliers[col], col] = ma[outliers[col]]
14
```

## 5. Normalize the data

```
125 def preprocess_entire_dataset(us_ds, mean, std):
126     exclude_columns = ['period_duration', 'region_type_id', 'table_id', 'property_type_id',
127                         'parent_metro_region_metro_code', 'year', 'month', 'price_increased_mom',
128                         'price_increased_yoy', 'fast_selling']
129
130     numeric_columns = us_ds.select_dtypes(include=['number']).columns
131
132     if exclude_columns:
133         columns_to_normalize = [col for col in numeric_columns if col not in exclude_columns]
134     else:
135         columns_to_normalize = numeric_columns
136
137     # normalize
138     print("Start normalizing...")
139     print("In preprocess entire...")
140     us_ds_normalized = us_ds.copy()
141     us_ds_normalized[columns_to_normalize] = normalize(us_ds[columns_to_normalize], mean, std)
```

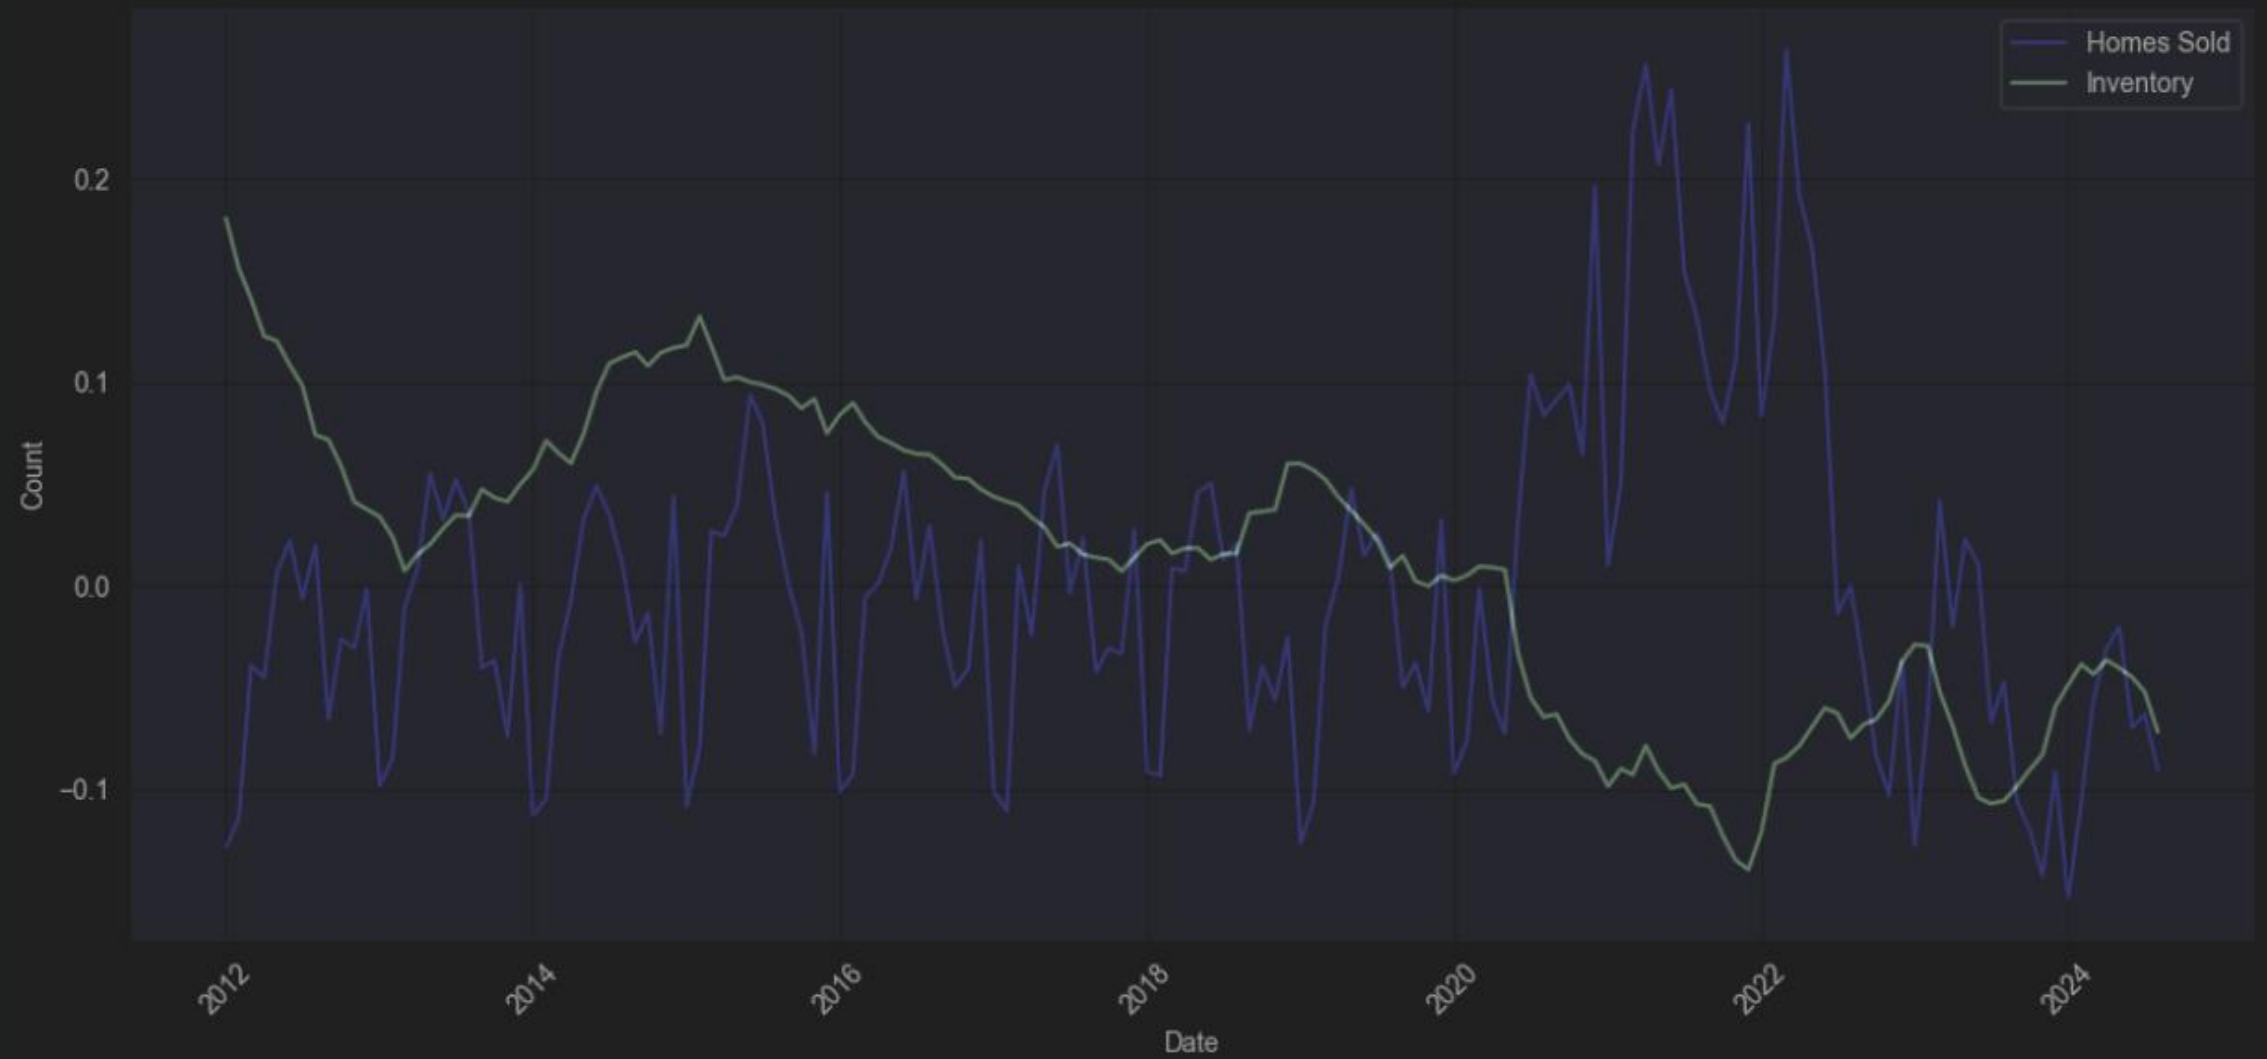
```
6 def normalize(us_ds, mean=None, std=None):
7     print("In normalize function...")
8     us_ds_normalized = us_ds.copy()
9
10     # Normalize only the selected numeric columns
11     us_ds_normalized = (us_ds_normalized - mean) / std
12
13     # us_ds_normalized = (us_ds - mean) / std
14     return us_ds_normalized
```

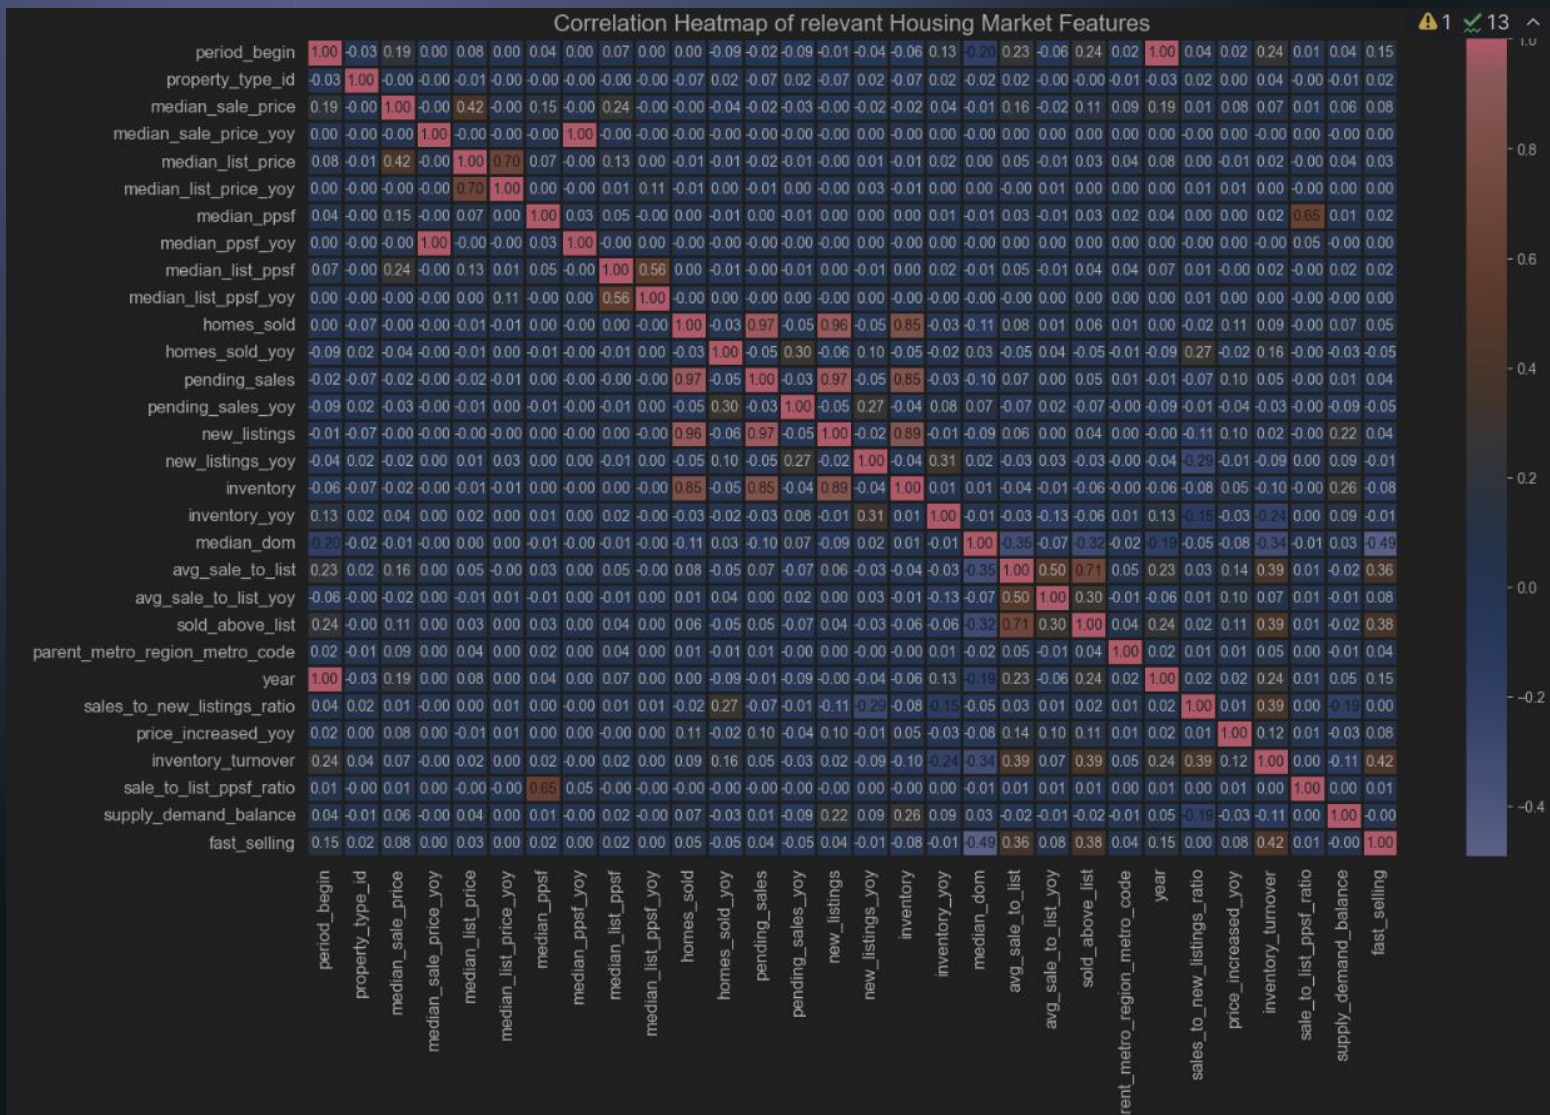

## 6. Generate diagrams





Trend of Homes Sold and Inventory Over Time (Monthly)





## 7. Encoding string columns

```
1 import category_encoders as ce
2
3 # Apply Binary Encoding
4 binary_encoder = ce.BinaryEncoder(cols=columns_for_encoding)
5 us_ds_binary = binary_encoder.fit_transform(us_ds)
6
```

state_code_0	state_code_1	state_code_2	state_code_3	state_code_4	state_code_5
0	0	0	0	0	1

property_type_0	property_type_1	property_type_2
0	0	1

## 8. Select the most relevant columns

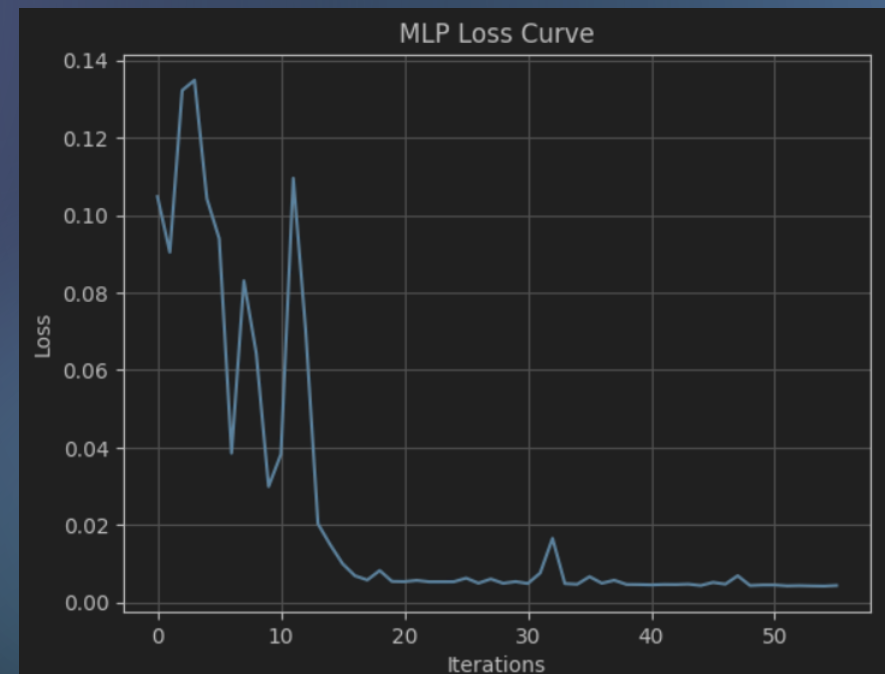
```
columns_for_model = [
    'state_code_0', 'state_code_1', 'state_code_2', 'state_code_3', 'state_code_4', 'state_code_5', 'property_type_0',
    'property_type_1', 'property_type_2', 'median_sale_price', 'median_dom', 'price_drops', 'inventory_turnover',
    'price_increased_mom', 'median_sale_price_mom', 'median_list_price', 'median_ppsf', 'median_list_ppsf',
    'avg_sale_to_list', 'sold_above_list', 'sale_to_list_ratio', 'sin_year', 'cos_year', 'sin_month', 'cos_month'
]
us_ds_for_model = us_ds[columns_for_model]
```

## 9. Model Training

```
1 from sklearn.model_selection import train_test_split
2 from sklearn.neural_network import MLPRegressor
3 from sklearn.metrics import mean_squared_error
4 [10]
5
6 y = us_ds['median_sale_price']
7
8 X = us_ds.drop(columns=['median_sale_price'])
9
10 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
11 [11]
12
13 mlp_model = MLPRegressor(hidden_layer_sizes=(64, 32), max_iter=200, random_state=42)
14 mlp_model.fit(X_train, y_train)
15 mlp_predictions = mlp_model.predict(X_test)
16 mlp_mse = mean_squared_error(y_test, mlp_predictions)
17 print(f"MLP Mean Squared Error: {mlp_mse}")
18 [13]
19
20 MLP Mean Squared Error: 0.012069494427731918
```

```
1 from sklearn.metrics import mean_absolute_error
2
3 mae = mean_absolute_error(y_test, mlp_predictions)
4 print(f"MAE: {mae}")
5 [41]
6
7 MAE: 0.022050769570323953
```

```
1 from sklearn.metrics import r2_score
2
3 r2 = r2_score(y_test, mlp_predictions)
4 print(f"R²: {r2}")
5 [42]
6
7 R²: 0.9875581700356271
```





# Bibliography

- [https://contrib.scikit-learn.org/category\\_encoders/](https://contrib.scikit-learn.org/category_encoders/)
- <https://developer.nvidia.com/blog/three-approaches-to-encoding-time-information-as-features-for-ml-models/>
- <https://www.datacamp.com/tutorial/multilayer-perceptrons-in-machine-learning>
- <https://towardsdatascience.com/multilayer-perceptron-explained-a-visual-guide-with-mini-2d-dataset-0ae8100c5d1c#:~:text=A%20Multilayer%20Perceptron%20%28MLP%29%20is%20a%20type%20of,connects%20to%20all%20nodes%20in%20the%20next%20layer.>
- <https://pythongeeks.org/data-preprocessing-in-machine-learning/>

Thank you!